

# COVID-19 Infection Analysis and Prediction

ROSINA Maxime

# Summary

<b>I - Analysis of the dataset</b>	<b>2</b>
Dataset comprehension and missing values	2
I.A - Correlation	4
I.B - PCA	5
<b>II - Bayes Nets</b>	<b>6</b>
II.A - Probability of having symptoms of Covid-19 after visiting Wuhan	6
II.B - Probability of being a true patient if having symptoms	6
II.C - Probability to die after visiting Wuhan	6
II.D - Average recovery day for a patient after visiting Wuhan	6
<b>III - Machine Learning</b>	<b>7</b>
III.A - K-Nearest Neighbours (K-NN) and Bayes Classification	7
III.B - Logistic Regression	7
III.C - K-Means	8
<b>IV - Improving the results and Theoretical formalism</b>	<b>8</b>
IV.A - Unbalanced data	9
IV.B - Manage missing values	9
IV.C - Grid-search algorithm	9
IV.D - Best model	9
<b>Conclusion</b>	<b>10</b>

# I - Analysis of the dataset

After numerous months affected by a sanitary crisis, the epidemiological situation is still evolving. I use this [dataset](#) which is daily updated. I downloaded my dataset December, 11, 2020.

The aim of this project is to apply artificial intelligence methods to this dataset in order to make analyzes and predictions to help the society to better understand the spread of the COVID-19 infection.

## Dataset comprehension and missing values

In order to understand all the attributes contained in this dataset, you can take a look at some scientific web papers such as Nature<sup>1</sup>. In this database, each row represents a single individual case.

To begin, take a look at some important information such as the shape of the dataset, the missing values, the outliers, and the correlation between attributes.

The dataset is composed of 33 columns and 2,676,311 rows. There are a lot of missing values.

Checking missing values			
ID	0	reported_market_exposure	2675242
age	2098293	additional_information	2630456
sex	2096154	chronic_disease_binary	0
city	977681	chronic_disease	2676096
province	452664	source	566964
country	115	sequence_available	2676299
latitude	61	outcome	2368929
longitude	61	date_death_or_discharge	2673163
geo_resolution	61	notes_for_discussion	2675671
date_onset_symptoms	2414712	location	2662935
date_admission_hospital	2560100	admin3	2595877
date_confirmation	108489	admin2	1850257
symptoms	2674259	admin1	1418753
lives_in_wuhan	2671973	country_new	30553
travel_history_dates	2673700	admin_id	61
travel_history_location	2667089	data_moderator_initials	933328
		travel_history_binary	65579

First I dealt with the missing values coming from the column age. I replaced all the missing values from age with the median age. I choose to remplace by the median because : *"If the variable is skewed, the mean is biased by the values at the far end of the*

<sup>1</sup> "Epidemiological data from the COVID-19 outbreak ... - Nature." 24 mars. 2020, <https://www.nature.com/articles/s41597-020-0448-0>. Date de consultation : 13 déc.. 2020.

*distribution. Therefore, the median is a better representation of the majority of the values in the variable.”<sup>2</sup>*

Then, I dropped all the rows with missing values from the attributes that I wanted to keep (sex, country, lives\_in\_Wuhan, chronic\_disease\_binary, travel\_history\_dates, outcome).

Finally, I had to manage a problem with the outcome column. In fact, data were not formalized, and many descriptions were given ('Recovered', 'recovered', 'recover' etc.). I formalized all the descriptions with three different values : 'died', 'hospitalized' and 'discharged'.

- ❑ 'died' is the value for people who die from COVID-19
- ❑ 'hospitalized' is the value for the people who are currently hospitalized because of COVID-19
- ❑ 'discharged' is the value for the people that are leaving the hospital (including all recovering people).

	age	sex	country	lives_in_Wuhan	chronic_disease_binary	outcome
504	35	0	23	0	0.0	0
756	35	0	23	1	0.0	0
765	35	0	23	0	0.0	0
797	4	0	13	0	0.0	0
798	39	0	13	0	0.0	1
...	...	...	...	...	...	...
637804	35	0	11	0	1.0	1
657538	35	1	11	0	0.0	0
658592	35	0	11	0	0.0	1
658603	35	1	11	0	0.0	1
672418	30	0	45	0	1.0	1


1583 rows x 6 columns

Our dataset is now ready !

<sup>2</sup> "Feature Engineering Part-1 Mean/ Median Imputation. | by ...." 17 août. 2020, <https://medium.com/analytics-vidhya/feature-engineering-part-1-mean-median-imputation-761043b95379>. Date de consultation : 13 déc.. 2020.

## I.A - Correlation

I took a look at the correlation between the different attributes.



I obtained the correlation matrix above. We can see that the most correlated variables are :


- ❑ lives\_in\_Wuhan and country with -20%
- ❑ age and chronic\_disease\_binary with 8%
- ❑ country and chronic\_disease\_binary with 5.6%
- ❑ outcome and chronic\_disease\_binary with 4.6%

Any of the correlations above are surprising. The country and living in Wuhan are the most correlated because a lot of patients that were tested were coming from Wuhan. Age and chronic disease are correlated because most chronic illnesses happen when you are old. Outcome and chronic\_disease\_binary are also correlated because we know,

nowadays, that an important part of the disease illness are coming from people who have some chronic illness in addition to the COVID-19.

## I.B - PCA

Now we take a closer look with a 2 dimensions PCA.



We can see that the data is very unbalanced between 'died', 'hospitalized' and 'discharged' values.

## II - Bayes Nets

We now compute some probability according to the data coming from our dataset.

### II.A - Probability of having symptoms of Covid-19 after visiting Wuhan

The probability for a person to have symptoms of COVID-19 if this person visited Wuhan is 60.63 %

### II.B - Probability of being a true patient if having symptoms

Probability for a person to be a true patient if this person has symptoms of COVID-19 and this person visited Wuhan is 98.42 %.

### II.C - Probability to die after visiting Wuhan

Probability for a person to die if this person visited Wuhan is 25.32%.


### II.D - Average recovery day for a patient after visiting Wuhan

Average recovery interval for a patient if this person visited Wuhan is 18 days.

## III - Machine Learning


### III.A - K-Nearest Neighbours (K-NN) and Bayes Classification

	precision	recall	f1-score	support
0	0.64	0.99	0.78	253
1	0.60	0.02	0.04	143
accuracy			0.64	396
macro avg	0.62	0.51	0.41	396
weighted avg	0.63	0.64	0.51	396



KNN model and confusion matrix

	precision	recall	f1-score	support
0	0.67	0.83	0.74	253
1	0.46	0.27	0.34	143
accuracy			0.62	396
macro avg	0.56	0.55	0.54	396
weighted avg	0.59	0.62	0.59	396



Naive Bayes model and confusion matrix

To have better results I think that we need to work with more balanced data, and with a larger sample. In fact, we only have 67% of precision in the best case that is not really high.

### III.B - Logistic Regression

Now I will use logistic regression in order to predict people age based on sex, country, lives\_in\_Wuhan, chronic\_disease\_binary, travel\_history\_dates, and outcome.

To do so I created some age categories in another column ageGroup. The four categories were :

- ❑ 0 for people between 0 and 20 years old
- ❑ 1 for people between 20 and 40 years old




- ❑ 2 for people between 40 and 60 years old
- ❑ 3 for people between 60 years old and more (up to 150 years old)

I applied Logistic Regression and I got the following results.

	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	12
1.0	0.36	0.39	0.38	145
2.0	0.41	0.58	0.48	166
3.0	0.20	0.01	0.03	73
accuracy			0.39	396
macro avg	0.24	0.25	0.22	396
weighted avg	0.34	0.39	0.34	396

MSE : 0.898989898989899



### Logistic regression model, MSE and confusion matrix

The following results are quite surprising because we don't have many people in category 0. The data is once again unbalanced. But we have a good MSE (0.989). We can clearly see that the logistic regression model better predicts people in age category 1 and 2 than category 0 and 3. However the precision is very low.


## III.C - K-Means

Now, we will apply K-Means to clusterize our data.

```
For n_clusters = 2 The average silhouette_score is : 0.42573633101231734
For n_clusters = 3 The average silhouette_score is : 0.41617989517390863
For n_clusters = 4 The average silhouette_score is : 0.415136807747351
For n_clusters = 5 The average silhouette_score is : 0.42879552421214434
For n_clusters = 6 The average silhouette_score is : 0.45315866065361193
```

From the above silhouette score we can see that the score is higher with 6 clusters.

**Silhouette analysis for KMeans clustering on sample data with n\_clusters = 6**



## IV - Improving the results and Theoretical formalism

### IV.A - Unbalanced data

The data is unbalanced, especially for the column outcome. We have :

hospitalized	997
discharged	558
died	28


If we reduce randomly the majority class the balance of the data will increase but it is not enough. In fact, we have only 28 persons who died in comparison to more than 500. The models will certainly increase in precision.

### IV.B - Manage missing values

We could manage the missing values in a better way instead of just deleting rows with missing values. For rows with binary values we could analyze and replace the missing values with the one with the highest probability.

### IV.C - Grid-search algorithm

The grid-search algorithm is an automation algorithm used for Hyperparameter optimization. Hyperparameter optimization is the problem of choosing an optimal set of optimal parameters to optimize a model.



After applying the grid search algorithm to our different models I obtained the table below :

Model	Naive Bayes classifier	KNN	Logistic regression
Best score	0.6621848739495798	0.6529269334852585	0.4035393818544367
Best(s) parameter(s)	var_smoothing = 1.0	n_neighbors = 18	C = 10.0 penalty = 'l2'

We can say that the best model in our case is the Naive Bayes classifier, with a var\_smoothing at 1.

#### IV.D - Best model



In our case, the best model is the Naive Bayes classifier. The parameter var\_smoothing gives the best results when it sets at 1. This parameter is the portion of the largest variance of all features that is added to variances for calculation stability. The second parameter that we have not studied here is the priors corresponding to the prior probabilities of the classes.

## Conclusion



We applied some of the machine learning models to our Covid19 dataset. The results are that our dataset is not optimized to perform high accurate predictions. In order to overcome the problem, we need to better balance our data, to better manage our missing values, and to use the best model with the best parameters, here the Naive Bayes classifier.

# Appendix



**Silhouette analysis for KMeans clustering on sample data with  $n\_clusters = 2$**



**Silhouette analysis for KMeans clustering on sample data with  $n\_clusters = 3$**



**Silhouette analysis for KMeans clustering on sample data with  $n\_clusters = 4$**



**Silhouette analysis for KMeans clustering on sample data with  $n\_clusters = 5$** 